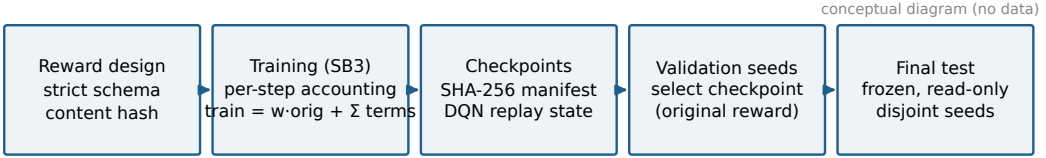
conceptual diagram (no data)
Reward design
strict schema
content hash
Training (SB3)
per-step accounting
train = w·orig + Σ terms
Checkpoints
SHA-256 manifest
DQN replay state
Validation seeds
select checkpoint
(original reward)
Final test
frozen, read-only
disjoint seeds
Audit outputs: original vs training return per term · reward-exploitation diagnostics · statistics across training seeds
per-step replay and re-scoring · provenance: configuration hash, seeds, versions, source revision, checkpoint digests